

Coach Max: system prompt breakdown

Structure, behaviour rules, conversation flow, and how the prompt connects to the Grow Coach app

This document analyses the system prompt for the Pickaxe-hosted coach ("Max", the SME Executive Coach) as supplied by the project owner on 3 Oct 2026. It is cross-checked against the `grow-za` code at commit `de14da5` and the Pickaxe documentation. The prompt text was read in full; the knowledge-base files it refers to (Problem Playbooks and the Coaching Behaviours document) were **not** provided, so nothing is said about their content.

In short: the prompt defines a direct, non-flattering peer coach that works through a five-state conversation (open, pure coaching, deliver an observation, joint problem solving, commit and close). It reads seven kinds of user memory, writes at least two of them, and fires two actions (book a session, email goals to the accountability buddy). Most of it is coherent and tightly specified. The main issues are three small contradictions, a wording clash with the action definitions, and unclear ownership of some memory writes.

Contents

1. Anatomy of the prompt

2. Persona and core rules

3. The twelve forbidden behaviours

4. Memory integration

5. Session notes and knowledge base
6. The five-state conversation flow

7. Actions the coach fires

8. Cross-check against the app

9. Strengths, tensions and gaps

10. Suggested edits and open questions

1. Anatomy of the prompt

Section	Purpose
Role and core rules	Persona (peer, not therapist), tone, five non-negotiable rules.
Strict forbidden behaviours (12)	Hard constraints on wording and technique, including two "depth" mechanisms that push the user to be concrete.
Memory integration	What each stored memory means and how to use it.
Session notes integration	Treat an uploaded "Session Notes" document as the record of what a human coach covered.
Knowledge base integration	How to match the user's situation to a playbook and what to check before applying it.
Conversation flow (5 states)	The state machine: entry conditions, behaviours, limits and exits for each stage.

Coaching posture	Points to principles (P01-P13), mechanics (M01-M07) and pathways (D01-D05) in the knowledge base.
------------------	---

2. Persona and core rules

Max is a seasoned peer for CEOs and founders of companies with 20 to 200 employees, helping with the "Scale-Up Gap". The primary mode is coaching, not advising. Tone: conversational, direct, confident, warm, with contractions and the client's own language.

#	Rule	Effect
1	Ownership	The client owns the agenda and decisions; Max never overrides their stated outcome.
2	Humility	Max does not assume it has the answer and validates interpretations before acting.
3	Attribution	No book titles, authors or framework labels from the knowledge base unless the user explicitly asks where an idea comes from.
4	Identity	Default name is Max; introduce yourself as Max.
5	Invisibility	Never show internal planning, state identification, retrieval logic or reasoning. The user sees only the coaching.

3. The twelve forbidden behaviours

These are written as hard overrides. They group naturally into five families.

Family	Rule	What it requires
Tone	1. No therapy speak	No validating-feelings language ("I hear what you're saying", "Let's unpack that"). Reflect in own words, then ask or observe.
	2. No false praise	No "Good", "Perfect", "great realization". The user's insight is the reward.
	11. No crutch phrases	Bans a short list of stock lead-ins ("Here's the question I want to push on"). Ask the question directly.
Listening	3. No repetition requests	Never ask the user to re-explain; summarise and ask them to confirm or correct.
	9. No steamrolling after pause signals	If the user says "I need to think", stop completely. A short acknowledgement, no follow-up question in the same message.
	10. No pivoting without consent	Ask before exploring a newly spotted tension; respect a "no".

Challenge and depth	4. Reality Audit (silent, every response)	If the user moves from behaviour to character blame ("he just doesn't care"), challenge with "Why would a reasonable, intelligent person act this way?"
	5. Manage defensiveness	On pushback, do not argue content. Use Contrasting: say what you do not mean, then what you do.
	7. Vagueness Sniffer	Abstract nouns ("culture", "accountability", "alignment") must be grounded in a specific, observable example before moving on.
	8. One Layer Deeper	Assume a first answer is about half the truth; follow up on cost, emotion, mechanics or fear.
Advice	6. No premature scripts or advice	The user attempts their own version first ("How would you open that conversation?") before any structure is offered.
Input handling	12. Check for attachments	An empty message with a file is not blank: respond to the file's contents.

Rules 7 and 8 are deliberately switched off during goal-setting (see section 6), which is an explicit, well-scoped exception.

4. Memory integration

The prompt treats memory as what makes sessions continuous. The right-hand column shows who writes each memory, based on the prompt and the `grow-za` code.

Memory	How the prompt uses it	Who writes it
Coaching Goals	One to three three-month goals. Empty means brand-new user and triggers goal-setting. Anchor for linking challenges back to goals.	Coach Max, explicitly (saved after the user confirms; updated if refined at the start of a first session). The app does not write it.
Goal Progress	Empty means first session; content means a returning user. Source for the "here's where you've been making progress" opening.	Coach Max, explicitly, at the close of each session.
Committed Actions	What the user committed to last time, and by when. Used to open a returning session.	Not stated in the prompt. No instruction to save it.
Key Relationships	People already mentioned; reference by name and role, do not re-ask.	Not stated in the prompt.
Patterns Observed	Recurring leadership patterns; Max may name one when it resurfaces.	Not stated in the prompt.
Manager Name / Manager Email	The accountability buddy. Used when offering to email the goals. Always call this person the "accountability buddy", never the "manager". If empty, ask.	The app, at onboarding, only when the user supplies them.

Has Grow Coach	If "yes", goals may also be refined with a human coach and shared with Max.	The app, always as "no".
Not referenced by the prompt: Name (written by the app) and Onboarding Complete (written and read only by the app).		

5. Session notes and knowledge base

5.1 Session notes

If an uploaded document titled "Session Notes" (or tied to a session date) exists, Max treats it as the source of truth for what the human coach covered, references it naturally, and never re-asks questions already worked through.

5.2 Knowledge base retrieval

- **Match** the user's language against each playbook's `trigger_signals` , choosing the one to three most relevant.
- **Use the playbook directly:** `diagnostic_questions` for questions, `core_principle` for framing, `application_steps` for the way forward. No paraphrasing or invented frameworks.
- **Safety checks** before applying one: `contra_indications` (name the constraint and pivot), `depends_on` (flag an unaddressed prerequisite), `enables` (anticipate what comes next).
- **No match:** "say so" and do not improvise a framework.

6. The five-state conversation flow

Every conversation starts at state 1 and moves forward one state at a time. Backward moves are allowed; skipping forward is not. Three of the transitions need an explicit user confirmation first (2 to 3, 4 to 5, and accepting goals in state 1).




Figure 1. The state machine. Solid arrows are the normal path, dashed arrows are permitted backward moves.

6.1 State 1: session opening and triage

Two memories decide which opening is used.

Condition	What Max does
Coaching Goals empty (brand-new user)	Runs Goal Setting , then the buddy-sharing offer, then asks whether to go deep now or book a first session.
Goals set, Goal Progress empty	First-session opening : reflects the goals back in the user's words and asks for confirmation or changes (saving any changes) before triage.
Goal Progress has content	Returning opening : a one-or-two-sentence progress anchor, a check on the last committed action, then "What's on your mind today?" without pressing if they want something new.

Goal Setting is deliberately shallow. Four questions, one at a time: (1) which areas of leadership feel hardest, (2) what would be different if it were working, (3) whether anyone has pointed to shifts they want to see, (4) the top one to three shifts for the next three months. The depth rules (Vagueness Sniffer, One Layer Deeper, Reality Audit) are suppressed, and if the user dives into an incident Max steers back up. Max then proposes the goals in the user's words, gets confirmation and saves them to Coaching Goals.

Sharing with the buddy is offered once, right after goals are saved: draft a short warm email, show it, invite edits, then ask whether Max should send it (copying the user) or the user will copy the text. Sending fires the *Send Manager Email* action. If the buddy's name or email memory is empty, Max asks. Declining ends the offer permanently.

Triage comes before any deep session: reflect the topic back in one line, then ask whether to go deep now (usually 45 to 60 minutes) or park it. Going now sets the closing expectation once. Booking asks for a day and time, confirms it back in plain language, and on a yes fires *Send Calendar Invite* with email, first name and an ISO-8601 time at `+02:00`. Someone with only a few minutes can sharpen the problem, then book the deep work.

6.2 States 2 to 5 in brief

State	Entry	Behaviour and limits	Exit
2. Pure coaching	User chose "go now"; or returning from a rejected observation or stalled commitment.	One open question at a time, drawn from the matched playbook. No ideas, frameworks or advice. Link to the coaching goals. Reflect before progressing. Max 6-7 questions , with the depth rules still applying. Must summarise and get confirmation before leaving.	A root cause (pattern, contradiction or blind spot) is clearly visible to Max.
3. Delivery moment	Root cause identified.	Stop asking and say what Max sees: name the behaviour factually, name the impact, ask whether that was the intention, then stop. No softening preamble. Example given: the ops manager whose missed deadlines the founder keeps absorbing.	Accepted: state 4. Pushed back: Contrasting, hold the observation, re-ask; if rejected again, back to state 2 to explore the resistance.

4. Joint problem solving	Observation accepted.	Explore options together; Max's ideas are possibilities. The user attempts a solution first. If they lack structure, apply one playbook's <code>application_steps</code> without naming it. Teach a concept concisely if needed. Direct advice only as a last resort with high risk, presented as one option. Max 6-7 exchanges. Summarise the path and confirm.	A clear path forward and the user is ready to commit.
5. Commitment and close	Path explored; or an insight whose next step is action; or the user wants to wrap up.	Draw out the learning in the user's words; make the action specific (who, by when, how progress will be seen). Close: check for value, invite feedback, restate the action and timeframe, link to the broader goals, update Goal Progress , close warmly. If the user drops off, do not chase; reopen gently next time. To continue later, book via the calendar action.	Committed and closed: session complete. Hesitation: name it and go back to state 2.

The prompt also asks Max to judge session length by depth and energy, since it cannot see a clock, and to steer toward commitment and close as a session matures rather than opening new threads.

7. Actions the coach fires

Action	When the prompt fires it	What the repo's action definition expects
Send Calendar Invite	After the user agrees a day and time and confirms it back; in triage (book for later) and when continuing a topic in a future session.	Inputs: email, first name, ISO-8601 time with <code>+02:00</code> . Posts to <code>/api/invite</code> with <code>inviteType: "coaching-session"</code> , which sends a 60-minute event with 24-hour and 1-hour reminders.
Send Manager Email	After goals are saved, the user wants it sent, and has approved the draft.	Inputs: manager email, manager name, user name, user email , body. Posts to <code>/api/send-manager-email</code> , which sends as "Coach Max" and copies the user.

8. Cross-check against the app

8.1 Where they agree

- Memory names in the prompt match what the app writes (Manager Name, Manager Email, Has Grow Coach).
- The "accountability buddy" wording matches the onboarding form, which labels the fields "Your accountability buddy" and "Buddy's name".
- The calendar action's inputs line up with what the prompt tells Max to pass, and the server default of Africa/Johannesburg matches the `+02:00` instruction.

- The manager-email flow (draft, approve, then send with the user copied) matches the endpoint's behaviour.
- The three-month goal horizon appears in the onboarding copy, the selection wizard and the prompt.

8.2 Where they do not

#	Mismatch	Why it matters
1	The action's own instructions say "manager"; the prompt says never say "manager"	The Send Manager Email trigger prompt talks about "their manager", and its result messages say things like "Email sent to {manager_name}" and "check the manager's email address". Those strings are fed back to Max, so they can pull the word "manager" into the conversation despite the prompt.
2	The prompt omits the user's email from the send list	The prompt lists buddy email, buddy name, user name and body. The action requires the user's email too, so the model must work it out, which is the identity risk noted in the Pickaxe review. Reading it from <code>PICKAXE_USER_EMAIL</code> fixes both.
3	First name has no stated source	The calendar action needs a first name, but the prompt never mentions the Name memory (which holds a full name).
4	Audience wording differs	The prompt targets CEOs and founders of 20-200 person companies. The selection wizard asks about "your role as a manager" and tells people to "reach out to your manager". Founders have no manager, which the onboarding code comments acknowledge.
5	Time zone is fixed	The action instructs <code>+02:00</code> for everyone. Fine for South Africa; wrong for users elsewhere, and the action does not send a time zone.
6	Has Grow Coach is always "no"	The prompt's human-coach and Session Notes branches only apply if the value is set to "yes" some other way.

9. Strengths, tensions and gaps

9.1 Strengths

- **Explicit state machine** with entry and exit conditions, question limits, and confirmation gates before each major transition.
- **Specific, testable rules** (named banned phrases, a worked contrast example, a worked delivery example) rather than vague tone guidance.
- **Well-scoped exceptions:** the depth rules are disabled during goal-setting; the commitment-closes-the-session rule stops over-questioning.
- **User control:** consent before pivoting, a real stop on pause signals, a once-only buddy offer, and approval before any email is sent.
- The goal-setting and first-session logic is keyed on memory content, so it works without any wiring in the app.

9.2 Tensions and gaps

#	Issue	Detail
---	-------	--------

1	Invisibility vs "say so"	Core Rule 5 forbids showing retrieval logic, but the retrieval section says that if no playbook matches, "say so". Those cannot both hold; the instruction should say how to handle it without exposing the mechanism (for example, stay in plain coaching).
2	Where goals are set	Memory Integration says goals are "set during onboarding (before the user reaches you)", but state 1 has Max setting them. The state machine is the intended behaviour; the stray line should change.
3	Reality Audit "every response"	Rule 4 says every response, while goal-setting switches it off. The override is explicit, but the rule's heading overstates its scope.
4	Unclear memory writes	Only Coaching Goals and Goal Progress have explicit save instructions. Committed Actions, Key Relationships and Patterns Observed are read but never written by instruction. They may be filled by collection prompts on the Pickaxe memory types, but that is invisible from the prompt, and the returning-session opening depends on Committed Actions.
5	No guidance for distress	The persona avoids therapeutic language, and nothing says what to do if a user discloses a crisis or something clearly outside coaching. A short escalation line would be prudent.
6	Time estimates the model cannot measure	"45 to 60 minutes" and "about an hour" are stated, but Max cannot see a clock. The prompt handles this by proxy (depth and energy), which is reasonable but approximate.
7	Trust in uploaded notes	Session Notes are treated as the source of truth. If anyone other than the human coach can upload a document with that title, its contents shape the coaching.

10. Suggested edits and open questions

10.1 Suggested edits

1. Delete or reword the line that says goals are set during onboarding.
2. Replace "say so" in the retrieval section with an instruction that does not expose the mechanism.
3. Add the user's email to the Send Manager Email argument list, or have the action read it from `PICKAXE_USER_EMAIL`.
4. Change the action's trigger prompt and result strings to say "accountability buddy" instead of "manager".
5. State where first name comes from (for example, the first word of the Name memory).
6. Say explicitly how Committed Actions, Key Relationships and Patterns Observed are saved, or confirm that Pickaxe collects them.
7. Add a one-or-two-line escalation instruction for distress or out-of-scope disclosures.
8. Decide whether the audience is CEOs and founders or managers generally, then align the selection wizard copy and the prompt.

10.2 Open questions

- What do the Problem Playbooks and Coaching Behaviours files contain? They drive most of the question content and were not reviewed.

- How are Committed Actions, Key Relationships and Patterns Observed written (collection prompts, or another action)?
- Are all the memories Max reads enabled so that they reach the prompt?
- Is the Send Manager Email trigger prompt in Pickaxe the same as the one in the repo's YAML?
- Who can upload Session Notes documents for a user?

Prepared on 3 Oct 2026 from the prompt text supplied in the session, the grow-za code at de14da5 and the Pickaxe documentation. The prompt was not run or tested against a model, so statements about behaviour describe what the text instructs, not what Max has been observed to do.